

Case Study: Statistical Test

Warren Fernandes 8940

Abhi Gupta 8944

Vinyas Kulal 8949

Liny Mathew 8950

In []:

```
import pandas as pd
import numpy as np
```

In [3]:

```
application_record = pd.read_csv('application_record.csv')
credit_record = pd.read_csv('credit_record.csv')
```


In [4]:

```
application_record
```

Out[4]:

	ID	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_TOTAL	NAME_INCOME_TYPE	NA
0	5008804	M	Y	Y	0	427500.0	Working	
1	5008805	M	Y	Y	0	427500.0	Working	
2	5008806	M	Y	Y	0	112500.0	Working	
3	5008808	F	N	Y	0	270000.0	Commercial associate	
4	5008809	F	N	Y	0	270000.0	Commercial associate	
...
438552	6840104	M	N	Y	0	135000.0	Pensioner	
438553	6840222	F	N	N	0	103500.0	Working	
438554	6841878	F	N	N	0	54000.0	Commercial associate	
438555	6842765	F	N	Y	0	72000.0	Pensioner	
438556	6842885	F	N	Y	0	121500.0	Working	

438557 rows × 18 columns



In [5]:

```
credit_record
```

Out[5]:

	ID	MONTHS_BALANCE	STATUS
0	5001711	0	X
1	5001711	-1	0
2	5001711	-2	0
3	5001711	-3	0
4	5001712	0	C
...
1048570	5150487	-25	C
1048571	5150487	-26	C
1048572	5150487	-27	C
1048573	5150487	-28	C
1048574	5150487	-29	C

1048575 rows × 3 columns

In [6]:

```
# Replace X,C values with 0 as they are identified as Good clients
credit_record.replace(['X','C'], 0,inplace=True)
```

In [7]:

```
credit_record.STATUS = pd.to_numeric(credit_record.STATUS)
```

In [8]:

```
# Searching for customers who have at Least one Late month
drop_ls = []
for i in range(len(credit_record)):
    if credit_record.STATUS[i] != 0:
        drop_ls.append(credit_record.ID[i])
```

In [9]:

```
len(drop_ls)
```

Out[9]:

14194

In [10]:

```
# Changing the STATUS of any client with at Least one Late month to 1
for i in range(len(credit_record)):
    if credit_record.ID[i] in drop_ls:
        credit_record.STATUS[i] = 1
```

In [11]:

```
credit_record.STATUS.value_counts()
```

Out[11]:

```
0    904764
1    143811
Name: STATUS, dtype: int64
```

In [12]:

```
credit_record.drop_duplicates(inplace=True)
credit_record
```

Out[12]:

	ID	MONTHS_BALANCE	STATUS
0	5001711	0	0
1	5001711	-1	0
2	5001711	-2	0
3	5001711	-3	0
4	5001712	0	0
...
1048570	5150487	-25	0
1048571	5150487	-26	0
1048572	5150487	-27	0
1048573	5150487	-28	0
1048574	5150487	-29	0

1048575 rows × 3 columns

In [13]:

```
print(f'No. of IDs in application_record = {len(application_record.ID)} No. of IDs in credit_record = {len(credit_record.ID)}')
```

No. of IDs in application_record = 438557 No. of IDs in credit_record = 1048575

In [14]:

```
dataset = application_record.merge(credit_record, on=['ID'], how='inner')
# on to choose which column to merger on
# How to get merge only the intersection between them
```

In [15]:

```
dataset.drop(['ID'], inplace=True, axis=1)
```

In [16]:

```
dataset.duplicated().sum()
```

Out[16]:

412393

In [17]:

```
dataset.drop_duplicates(inplace=True)
```

In [18]:

dataset

Out[18]:

	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_TOTAL	NAME_INCOME_TYPE	NAME_EDUC
0	M	Y	Y	0	427500.0	Working	Hig
1	M	Y	Y	0	427500.0	Working	Hig
2	M	Y	Y	0	427500.0	Working	Hig
3	M	Y	Y	0	427500.0	Working	Hig
4	M	Y	Y	0	427500.0	Working	Hig
...
777710	M	N	Y	0	112500.0	Working	Seconda
777711	M	N	Y	0	112500.0	Working	Seconda
777712	M	N	Y	0	112500.0	Working	Seconda
777713	M	N	Y	0	112500.0	Working	Seconda
777714	M	N	Y	0	112500.0	Working	Seconda

365322 rows × 19 columns

In [19]:

dataset.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 365322 entries, 0 to 777714
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CODE_GENDER                          365322 non-null object
1   FLAG_OWN_CAR                         365322 non-null object
2   FLAG_OWN_REALTY                     365322 non-null object
3   CNT_CHILDREN                        365322 non-null int64
4   AMT_INCOME_TOTAL                   365322 non-null float64
5   NAME_INCOME_TYPE                   365322 non-null object
6   NAME_EDUCATION_TYPE                365322 non-null object
7   NAME_FAMILY_STATUS                 365322 non-null object
8   NAME_HOUSING_TYPE                  365322 non-null object
9   DAYS_BIRTH                         365322 non-null int64
10  DAYS_EMPLOYED                      365322 non-null int64
11  FLAG_MOBIL                         365322 non-null int64
12  FLAG_WORK_PHONE                    365322 non-null int64
13  FLAG_PHONE                         365322 non-null int64
14  FLAG_EMAIL                         365322 non-null int64
15  OCCUPATION_TYPE                    252192 non-null object
16  CNT_FAM_MEMBERS                    365322 non-null float64
17  MONTHS_BALANCE                     365322 non-null int64
18  STATUS                             365322 non-null int64
dtypes: float64(2), int64(9), object(8)
memory usage: 55.7+ MB
```

In [20]:

```
dataset.describe()
```

Out[20]:

	CNT_CHILDREN	AMT_INCOME_TOTAL	DAYS_BIRTH	DAYS_EMPLOYED	FLAG_MOBIL	FLAG_WORK_PHONE	FLAG_PHONE	FLAG
count	365322.000000	3.653220e+05	365322.000000	365322.000000	365322.0	365322.000000	365322.000000	365322
mean	0.425742	1.848982e+05	-16161.482656	60776.306365	1.0	0.221878	0.294214	0
std	0.768540	1.017316e+05	4144.182785	139028.719425	0.0	0.415510	0.455689	0
min	0.000000	2.700000e+04	-25152.000000	-15713.000000	1.0	0.000000	0.000000	0
25%	0.000000	1.170000e+05	-19614.000000	-3208.000000	1.0	0.000000	0.000000	0
50%	0.000000	1.575000e+05	-15849.000000	-1566.000000	1.0	0.000000	0.000000	0
75%	1.000000	2.250000e+05	-12676.000000	-378.000000	1.0	0.000000	1.000000	0
max	19.000000	1.575000e+06	-7489.000000	365243.000000	1.0	1.000000	1.000000	1

In [21]:

```
dataset.isna().sum()
```

Out[21]:

CODE_GENDER	0
FLAG_OWN_CAR	0
FLAG_OWN_REALTY	0
CNT_CHILDREN	0
AMT_INCOME_TOTAL	0
NAME_INCOME_TYPE	0
NAME_EDUCATION_TYPE	0
NAME_FAMILY_STATUS	0
NAME_HOUSING_TYPE	0
DAYS_BIRTH	0
DAYS_EMPLOYED	0
FLAG_MOBIL	0
FLAG_WORK_PHONE	0
FLAG_PHONE	0
FLAG_EMAIL	0
OCCUPATION_TYPE	113130
CNT_FAM_MEMBERS	0
MONTHS_BALANCE	0
STATUS	0

dtype: int64

In [22]:

```
dataset.isna().sum().sum()
```

Out[22]:

113130

In [23]:

```
dataset.OCCUPATION_TYPE
```

Out[23]:

0	NaN
1	NaN
2	NaN
3	NaN
4	NaN
...	
777710	Laborers
777711	Laborers
777712	Laborers
777713	Laborers
777714	Laborers

Name: OCCUPATION_TYPE, Length: 365322, dtype: object

In [24]:

```
dataset.OCCUPATION_TYPE.value_counts()
```

Out[24]:

```
Laborers          62839
Core staff        34175
Sales staff       33786
Managers          31066
Drivers           23349
High skill tech staff 14459
Medicine staff    11937
Accountants       11926
Security staff    6851
Cooking staff     6663
Cleaning staff    5201
Private service staff 2989
Low-skill Laborers 2000
Secretaries       1523
Waiters/barmen staff 1272
HR staff          973
IT staff          617
Realty agents     566
Name: OCCUPATION_TYPE, dtype: int64
```

In [25]:

```
dataset.OCCUPATION_TYPE.replace(np.nan, 'Other', inplace = True)
```

In [26]:

```
dataset.OCCUPATION_TYPE.value_counts()
```

Out[26]:

```
Other          113130
Laborers       62839
Core staff     34175
Sales staff    33786
Managers       31066
Drivers        23349
High skill tech staff 14459
Medicine staff 11937
Accountants    11926
Security staff 6851
Cooking staff  6663
Cleaning staff 5201
Private service staff 2989
Low-skill Laborers 2000
Secretaries    1523
Waiters/barmen staff 1272
HR staff       973
IT staff       617
Realty agents  566
Name: OCCUPATION_TYPE, dtype: int64
```

In [28]:

```
dataset
```

Out[28]:

	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_TOTAL	NAME_INCOME_TYPE	NAME_EDUC
0	M	Y	Y	0	427500.0	Working	Hig
1	M	Y	Y	0	427500.0	Working	Hig
2	M	Y	Y	0	427500.0	Working	Hig
3	M	Y	Y	0	427500.0	Working	Hig
4	M	Y	Y	0	427500.0	Working	Hig
...
777710	M	N	Y	0	112500.0	Working	Seconda
777711	M	N	Y	0	112500.0	Working	Seconda
777712	M	N	Y	0	112500.0	Working	Seconda
777713	M	N	Y	0	112500.0	Working	Seconda
777714	M	N	Y	0	112500.0	Working	Seconda

365322 rows × 19 columns

In []:

```
dataset.to_csv('credit.csv')
```

One Sample

Test for Proportion

In [89]:

```
sample = dataset.sample(frac=0.10)
```

In [91]:

```
sample.describe()
```

Out[91]:

	CNT_CHILDREN	AMT_INCOME_TOTAL	DAYS_BIRTH	DAYS_EMPLOYED	FLAG_MOBIL	FLAG_WORK_PHONE	FLAG_PHONE	FLAG_I
count	36532.000000	3.653200e+04	36532.000000	36532.000000	36532.0	36532.000000	36532.000000	36532.0
mean	0.429131	1.844514e+05	-16186.627724	61935.026443	1.0	0.221422	0.294810	0.0
std	0.769898	1.019580e+05	4154.943376	140002.017208	0.0	0.415210	0.455964	0.2
min	0.000000	2.700000e+04	-25152.000000	-15713.000000	1.0	0.000000	0.000000	0.0
25%	0.000000	1.170000e+05	-19661.000000	-3174.000000	1.0	0.000000	0.000000	0.0
50%	0.000000	1.575000e+05	-15849.000000	-1555.000000	1.0	0.000000	0.000000	0.0
75%	1.000000	2.250000e+05	-12705.000000	-356.000000	1.0	0.000000	1.000000	0.0
max	19.000000	1.575000e+06	-7757.000000	365243.000000	1.0	1.000000	1.000000	1.0

In [93]:

```
sample.CODE_GENDER.value_counts()
```

Out[93]:

```
F    23937
M    12595
Name: CODE_GENDER, dtype: int64
```

In [101]:

```
proportions_ztest(count=23937, nobs=36532, value=0.66)
```

Out[101]:

```
(-1.9166880588315112, 0.05527757319676388)
```

In [102]:

```
dataset['CODE_GENDER'].value_counts(normalize=True)
```

Out[102]:

```
F    0.655737  
M    0.344263  
Name: CODE_GENDER, dtype: float64
```

Test for mean

In [103]:

```
from statsmodels.stats.weightstats import ztest
```

In [124]:

```
ztest(x1=sample['AMT_INCOME_TOTAL'], value=190000)
```

Out[124]:

```
(-10.401568432043778, 2.4388536106918665e-25)
```

In [125]:

```
dataset['AMT_INCOME_TOTAL'].mean()
```

Out[125]:

```
184898.23890157178
```

Two Sample

In [156]:

```
sample1 = dataset.sample(frac=0.10)  
sample2 = dataset.sample(frac=0.10)
```



In [157]:

```
sample1
```

Out[157]:

	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_TOTAL	NAME_INCOME_TYPE	NAME_EDUC
183834	M	N	Y	0	135000.0	Working	Secondar
98641	M	Y	N	1	157500.0	Commercial associate	Secondar
573812	M	Y	Y	0	157500.0	Working	Secondar
577461	F	Y	Y	2	90000.0	Commercial associate	Secondar
383968	F	Y	Y	0	337500.0	Working	Secondar
...
381732	F	N	Y	0	90000.0	Working	Incc
34201	F	Y	Y	0	112500.0	Pensioner	Hig
515880	F	N	Y	0	72000.0	Pensioner	Secondar
211624	M	N	Y	0	112500.0	Commercial associate	Secondar
294057	M	N	N	3	58500.0	Working	Secondar

36532 rows × 19 columns




In [158]:

```
sample2
```

Out[158]:

	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_TOTAL	NAME_INCOME_TYPE	NAME_EDUCATIO
1281	M	Y	Y	2	450000.0	Working	Higher e
1724	M	N	N	0	135000.0	Commercial associate	Secondary / se
1776	M	Y	N	1	450000.0	Working	Higher e
1221	F	N	Y	0	135000.0	Working	Secondary / se
1349	M	Y	Y	0	270000.0	Working	Secondary / se
...
1323	F	N	N	0	180000.0	Working	Higher e
1185	F	N	Y	1	117000.0	Working	Secondary / se
1643	F	N	Y	3	67500.0	Commercial associate	Secondary / se
1972	M	N	Y	2	207000.0	Commercial associate	Secondary / se
1662	M	Y	Y	0	540000.0	Working	Secondary / se

32 rows × 19 columns



In [159]:

```
sample1['CODE_GENDER'].value_counts()
```

Out[159]:

F 23976
M 12556
Name: CODE_GENDER, dtype: int64

In [160]:

```
sample2['CODE_GENDER'].value_counts()
```

Out[160]:

```
F    23955
M    12577
Name: CODE_GENDER, dtype: int64
```

Difference of Two Proportions

In [161]:

```
count = np.array([23976, 23955])
nobs = np.array([36532, 36532])
proportions_ztest(count, nobs)
```

Out[161]:

```
(0.16354625601354528, 0.8700883563502826)
```

In [162]:

```
sample1['CODE_GENDER'].value_counts(normalize=True)
```

Out[162]:

```
F    0.656301
M    0.343699
Name: CODE_GENDER, dtype: float64
```

In [163]:

```
sample2['CODE_GENDER'].value_counts(normalize=True)
```

Out[163]:

```
F    0.655726
M    0.344274
Name: CODE_GENDER, dtype: float64
```

In [165]:

```
sample1['FLAG_OWN_CAR'].value_counts()
```

Out[165]:

```
N    22717
Y    13815
Name: FLAG_OWN_CAR, dtype: int64
```

In [166]:

```
sample2['FLAG_OWN_CAR'].value_counts()
```

Out[166]:

```
N    22526
Y    14006
Name: FLAG_OWN_CAR, dtype: int64
```

In [168]:

```
count = np.array([22717, 22526])
nobs = np.array([36532, 36532])
proportions_ztest(count, nobs, value=0.01)
```

Out[168]:

```
(-1.3281186463712324, 0.18413891442018981)
```

Difference of Two means

In [169]:

```
ztest(x1=sample1['AMT_INCOME_TOTAL'], x2=sample2['AMT_INCOME_TOTAL'], value=0)
```

Out[169]:

```
(-1.2696213576394317, 0.2042195378968562)
```

In [170]:

```
sample1['AMT_INCOME_TOTAL'].mean()
```

Out[170]:

```
184719.69859301436
```

In [171]:

```
sample2['AMT_INCOME_TOTAL'].mean()
```

Out[171]:

```
185672.07594711485
```

One Sample, Two Measures

Chisquare Test for independence

In [172]:

```
from statsmodels.stats.proportion import proportions_chisquare
```

In [176]:

```
np.array([23976, 23955])  
nobs = np.array([36532, 36532])  
proportions_chisquare(count, nobs)
```

Out[176]:

```
(2.117610336688899,  
 0.1456135874098347,  
 (array([[22717., 13815.],  
         [22526., 14006.]]),  
  array([[22621.5, 13910.5],  
         [22621.5, 13910.5]])))
```

In [179]:

```
count = np.array([22717, 22526])  
nobs = np.array([36532, 36532])  
proportions_chisquare(count, nobs)
```

Out[179]:

```
(2.117610336688899,  
 0.1456135874098347,  
 (array([[22717., 13815.],  
         [22526., 14006.]]),  
  array([[22621.5, 13910.5],  
         [22621.5, 13910.5]])))
```

Regression Analysis

In [180]:

```
from statsmodels.api import OLS
```

In [181]:

```
mod = OLS(dataset['DAYS_BIRTH'], dataset['AMT_INCOME_TOTAL'])
```

In [182]:

```
res = mod.fit()
```

In [183]:

```
print(res.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	DAYS_BIRTH	R-squared (uncentered):				0.703
Model:	OLS	Adj. R-squared (uncentered):				0.703
Method:	Least Squares	F-statistic:				8.668e+05
Date:	Tue, 07 Mar 2023	Prob (F-statistic):				0.00
Time:	18:02:42	Log-Likelihood:				-3.8481e+06
No. Observations:	365322	AIC:				7.696e+06
Df Residuals:	365321	BIC:				7.696e+06
Df Model:	1					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

AMT_INCOME_TOTAL	-0.0663	7.12e-05	-931.013	0.000	-0.066	-0.066
=====						
Omnibus:	132324.664	Durbin-Watson:			0.042	
Prob(Omnibus):	0.000	Jarque-Bera (JB):			1329668.249	
Skew:	1.457	Prob(JB):			0.00	
Kurtosis:	11.881	Cond. No.			1.00	
=====						

Notes:
[1] R² is computed without centering (uncentered) since the model does not contain a constant.
[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Difference of means

In [211]:

```
ztest(x1=sample1['CNT_FAM_MEMBERS'], x2=sample2['CNT_CHILDREN'], value=1.8)
```

Out[211]:

(-5.180909056271185, 2.208071348369464e-07)

In []: